

# Introduction - Machine Learning in Azure

April 14, 2018



Lending Club is a peer-to-peer lending company to connect borrowers (people who need money) with investors (people who have money).

This notebook will demonstrate following tasks:

- Build a machine learning model to predict whether or not borrower paid back their loan in full
- Operationalize the model by deploying it as a web service in Microsoft Azure platform
- Data can be downloaded in csv format from here <https://www.lendingclub.com/info/download-data.action> (<https://www.lendingclub.com/info/download-data.action>)
- This is a Python 2 notebook. Although Azure supports Python 3, its web services deployment does not support Python 3 coding yet

## ★ Azure Work Space Settings

Log in to Microsoft Azure Machine Learning Studio and, under SETTINGS, retrieve **WORKSPACE ID** and **AUTHORIZATION TOKEN**

In [1]:

```
# install libraries to local machine if needed
# pip install azure
# pip install azure-ml-api-sdk
# pip install azureml
```

In [68]:

```
import azureml

workspace_id = "enter workspace id from azure"
authorization_token = "enter authorization token from azure"
ws = azureml.Workspace(workspace_id=workspace_id, authorization_token=authorization_
```

## ★ Import Libraries and Dataset

In [69]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

# dataset was uploaded to Azure. import dataset from Azure
ds = ws.datasets["loan_data.csv"]

# convert dataset to pandas dataframe
df = ds.to_dataframe()
```

## ★ Explore Dataset and Perform Data Preprocessing

### § Check dataframe information

- "purpose" is a categorical feature
- "not.fully.paid" is the label/target variable

In [70]:

```
# check dataframe information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
credit.policy      9578 non-null int64
purpose           9578 non-null object
int.rate          9578 non-null float64
installment       9578 non-null float64
log.annual.inc    9578 non-null float64
dti               9578 non-null float64
fico              9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal         9578 non-null int64
revol.util        9578 non-null float64
inq.last.6mths    9578 non-null int64
delinq.2yrs       9578 non-null int64
pub.rec           9578 non-null int64
not.fully.paid    9578 non-null int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

### § View first five rows of data

In [71]:

```
df.head()
```

Out[71]:

	credit.policy		purpose	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.lin
0	1	debt_consolidation		0.1189	829.10	11.350407	19.48	737	5639.95833
1	1	credit_card		0.1071	228.22	11.082143	14.29	707	2760.00000
2	1	debt_consolidation		0.1357	366.86	10.373491	11.63	682	4710.00000
3	1	debt_consolidation		0.1008	162.34	11.350407	8.10	712	2699.95833
4	1	credit_card		0.1426	102.92	11.299732	14.97	667	4066.00000

**§ Initially columns use dot notation. We want to rename them to prevent future issues when creating web services**

In [72]:

```
df.rename( columns=(lambda value: value.replace(".", "_")), inplace=True )
df.head()
```

Out[72]:

	credit_policy		purpose	int_rate	installment	log_annual_inc	dti	fico	days_with_cr
0	1	debt_consolidation		0.1189	829.10	11.350407	19.48	737	5639.95
1	1	credit_card		0.1071	228.22	11.082143	14.29	707	2760.00
2	1	debt_consolidation		0.1357	366.86	10.373491	11.63	682	4710.00
3	1	debt_consolidation		0.1008	162.34	11.350407	8.10	712	2699.95
4	1	credit_card		0.1426	102.92	11.299732	14.97	667	4066.00

**§ Check if any feature has NULL value**

In [73]:

```
# check null value
df.isnull().sum()
```

Out[73]:

```
credit_policy      0
purpose            0
int_rate           0
installment        0
log_annual_inc     0
dti                0
fico               0
days_with_cr_line 0
revol_bal          0
revol_util         0
inq_last_6mths     0
delinq_2yrs        0
pub_rec            0
not_fully_paid     0
dtype: int64
```

## § View distribution of target variable

In [74]:

```
df["not_fully_paid"].value_counts()
```

Out[74]:

```
0      8045
1      1533
Name: not_fully_paid, dtype: int64
```

**§ Since "purpose" column is categorical, we want to transform it using one-hot-encoding, And remove one column to avoid multicollinearity.**

In [75]:

```
# show 7 purpose values
print( "There are 7 purpose values:\n{}".format(df["purpose"].value_counts()) )

df_final = pd.get_dummies(data=df, columns=["purpose"], drop_first=True)
df_final.filter(regex=("purpose.*")).head() # show one less purpose value (6)
```

There are 7 purpose values:

debt_consolidation	3957
all_other	2331
credit_card	1262
home_improvement	629
small_business	619
major_purchase	437
educational	343

Name: purpose, dtype: int64

Out[75]:

	purpose_credit_card	purpose_debt_consolidation	purpose_educational	purpose_home_improver
0	0	1	0	
1	1	0	0	
2	0	1	0	
3	0	1	0	
4	1	0	0	

In [76]:

```
# confirm final dataset information
df_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 19 columns):
credit_policy          9578 non-null int64
int_rate               9578 non-null float64
installment            9578 non-null float64
log_annual_inc         9578 non-null float64
dti                    9578 non-null float64
fico                   9578 non-null int64
days_with_cr_line     9578 non-null float64
revol_bal              9578 non-null int64
revol_util             9578 non-null float64
inq_last_6mths         9578 non-null int64
delinq_2yrs            9578 non-null int64
pub_rec                9578 non-null int64
not_fully_paid         9578 non-null int64
purpose_credit_card    9578 non-null uint8
purpose_debt_consolidation 9578 non-null uint8
purpose_educational    9578 non-null uint8
purpose_home_improvement 9578 non-null uint8
purpose_major_purchase 9578 non-null uint8
purpose_small_business 9578 non-null uint8
dtypes: float64(6), int64(7), uint8(6)
memory usage: 1.0 MB
```

## ★ Split data into training and test set

There are syntax differences between Python 3 and Python 2 for calling sklearn. Use following to check current running Python version:

```
import sys
print(sys.executable)
print(sys.version)
print(sys.version_info)
```

In [78]:

```
import platform; print(platform.platform())
import sys; print("Python", sys.version)
import numpy; print("NumPy", numpy.__version__)
import scipy; print("SciPy", scipy.__version__)
```

```
Darwin-16.7.0-x86_64-i386-64bit
('Python', '2.7.14 |Anaconda, Inc.| (default, Mar 27 2018, 12:28:59) \
n[GCC 4.2.1 Compatible Clang 4.0.1 (tags/RELEASE_401/final)]')
('NumPy', '1.14.2')
('SciPy', '1.0.1')
```

In [79]:

```
# python 3 version
# from sklearn.model_selection import train_test_split

# python 2 version
from sklearn.cross_validation import train_test_split

y = df_final["not_fully_paid"]
X = df_final.drop("not_fully_paid", axis=1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
                                                    random_state=42)
```

## ★ Build a Decision Tree model

In [80]:

```
from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
```

Out[80]:

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
None,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort=False, random_state=
None,
                        splitter='best')
```

## § Make predictions based on DT

In [81]:

```
# prediction based on decision tree  
dt_prediction = dt.predict(X_test)
```

## ★ Build a Random Forest model

In [82]:

```
from sklearn.ensemble import RandomForestClassifier  
  
rf = RandomForestClassifier()  
rf.fit(X_train, y_train)
```

Out[82]:

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                        max_depth=None, max_features='auto', max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,  
                        oob_score=False, random_state=None, verbose=0,  
                        warm_start=False)
```

## § Make predictions based on RF

In [83]:

```
rf_prediction = rf.predict(X_test)
```

## ★ Evaluate and compare model performance

Clearly Random Forest yields better result than single Decision Tree model



In [84]:

```
from sklearn.metrics import confusion_matrix, classification_report

print("*** From Decision Tree:")
print( confusion_matrix(y_test, dt_prediction) )
print( classification_report(y_test, dt_prediction) )

print("*** From Random Forest:")
print( confusion_matrix(y_test, rf_prediction) )
print( classification_report(y_test, rf_prediction) )
```

\*\*\* From Decision Tree:

```
[[2231  419]
 [ 399  112]]
```

	precision	recall	f1-score	support
0	0.85	0.84	0.85	2650
1	0.21	0.22	0.21	511
avg / total	0.75	0.74	0.74	3161

\*\*\* From Random Forest:

```
[[2603   47]
 [ 494   17]]
```

	precision	recall	f1-score	support
0	0.84	0.98	0.91	2650
1	0.27	0.03	0.06	511
avg / total	0.75	0.83	0.77	3161

## § Check and visualize model features importance

We'll choose Random Forest model

In [85]:


```
import seaborn as sns

# get feature importance
importances = rf.feature_importances_

# create a df containing column names and corresponding importances
list_feature = [pd.DataFrame(X.columns.values, columns=["names"]), pd.DataFrame(importances, columns=["importances"])
df_feature = pd.concat(list_feature, axis=1)

# sort in descending order before plotting
df_feature.sort_values(by="importances", ascending=False, inplace=True)

# plotting
plt.figure(figsize=(10,5))
g = sns.barplot(data=df_feature, x="names", y="importances")
g.set_xticklabels(g.get_xticklabels(), rotation=90)
plt.show()
```




In [86]:

```
# get feature importance
importances = rf.feature_importances_

# sort feature importances in descending order
indices = np.argsort(importances)[::-1]

# use list comprehension to extract list of feature names which has same order of i
names = [X.columns[i] for i in indices]

# plot impoirtnce in descending order
plt.figure(figsize=(10,5))
plt.bar(range(X.shape[1]), importances[indices])
plt.xticks(range(X.shape[1]), names, rotation=90)
plt.show()
```



★ Rebuild model using top 5 features (weighted by feature importances found earlier)

We'll use this simpler feature set to deploy web services

In [87]:

```
df_feature.iloc[:5,:]
```

Out[87]:

	names	importances
6	days_with_cr_line	0.115676
2	installment	0.114646
3	log_annual_inc	0.113702
7	revol_bal	0.108914
8	revol_util	0.108386

In [88]:

```
y_op = y

X_op = X[["days_with_cr_line", "log_annual_inc", "installment", "dti", "revol_util"]]
X_op_train, X_op_test, y_op_train, y_op_test = train_test_split(X_op, y_op, test_size=0.2)

# build a Random Forest model
rf.fit(X_op_train, y_op_train)

# prediction
prediction_op = rf.predict(X_op_test)

# evaluate
print( confusion_matrix(y_op_test, prediction_op) )
print( classification_report(y_op_test, prediction_op) )
```

[[2603 47]					
[ 500 11]]					
		precision	recall	f1-score	support
	0	0.84	0.98	0.90	2650
	1	0.19	0.02	0.04	511
avg / total		0.73	0.83	0.76	3161

## \$ Deploying the model as a web service

We create a wrapper function that takes input features as an argument. Then the function calls `predict()` method of our trained model and returns prediction result - classification as 0 (not fully paid) or 1 (fully paid)

## \$ Template for building web services

```
from azureml import services
@services.publish(..)
@services.types(..)
@services.returns(..)
def user_defined_function(..)
```

In [89]:

```
X_op.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 5 columns):
days_with_cr_line      9578 non-null float64
log_annual_inc          9578 non-null float64
installment             9578 non-null float64
dti                     9578 non-null float64
revol_util              9578 non-null float64
dtypes: float64(5)
memory usage: 374.2 KB
```

In [90]:

```
from azureml import services

# workspace_id and authorization_token are defined in the beginning of notebook
@services.publish(workspace_id, authorization_token)
# top 5 features with corresponding types
@services.types(days_with_cr_line=float, log_annual_inc=float, installment=float, dti=float, revol_util=float)
# predicted label
@services.returns(int)

# actual function called by web services
def credit_check(days_with_cr_line, log_annual_inc, installment, dti, revol_util):
    return rf.predict([days_with_cr_line, log_annual_inc, installment, dti, revol_util])

# for testing run within the notebook
service_url = credit_check.service.url
api_key = credit_check.service.api_key
help_url = credit_check.service.help_url
service_id = credit_check.service.service_id
```

## ★ Example for consuming web services from the notebook

In [91]:

```
# samples
# df[df["not_fully_paid"]==0][["days_with_cr_line", "log_annual_inc", "installment", "cred
# df[df["not_fully_paid"]==1][["days_with_cr_line", "log_annual_inc", "installment", "c

# test output
result = credit_check.service(2760.000000, 11.082143, 228.22, 14.29, 76.7)
# result = credit_check.service(4209.95, 11.884489, 678.08, 10.15, 74.1)

# print customized result
transform_result = lambda x: "NO" if x=="0" else "YES"
print( "fully paid? {}".format(transform_result(result)) )

# alternative approach to print result
# print( lambda x: "NO" if x==0 else "YES" )(result)
```

fully paid? YES

## ★ Example for consuming web services from an application

In [94]:

```
import urllib2
import json

example_values = [10.71, 4, 3180.041667, 76.8, 194.02]
# example_values = [4209.95, 11.884489, 678.08, 10.15, 74.1]

data = {"Inputs": {"input1": { "ColumnNames": ["days_with_cr_line", "log_annual_inc",
                                              "Values": [example_values] } }, # specified feature v
        "GlobalParameters": {} }

body = json.dumps(data)
headers = {'Content-Type': 'application/json', 'Authorization': ('Bearer ' + api_key)}
req = urllib2.Request(service_url, body, headers)

try:
    response = urllib2.urlopen(req)
    result = json.loads(response.read()) # load json-formatted string response as c
    text = result['Results']['output1']['value']['Values'][0][0] # convert numeric l
    transform_result = lambda x: "NO" if x=="0" else "YES"
    print("fully paid? {}={}".format(text, transform_result(text))) # get the returne
#     print(result['Results']['output1']['value']['Values'][0][0]) # get the returne
except urllib2.HTTPError, error:
    print("The request failed with status code: " + str(error.code))
    print(error.info())
    print(json.loads(error.read()))
```

fully paid? 0=NO

In [ ]: